

# Hadoop Test : 2017

## Test ka Aarambh

Choose the correct statement

1 point

- ☐ A) Datanode manage file system namespace
- ☒ B) Namenode stores metadata
- ☐ C) NameNode stores actual data
- ☐ D) None of the above

Point out the correct statement :

1 point

- ☐ a) HDFS provides low latency access to single rows from billions of records (Random access)
- ☒ b) HBase sits on top of the Hadoop File System and provides read and write access
- ☐ c) HBase is a distributed file system suitable for storing large files
- ☐ d) None of the mentioned

What data, does a Reducer reduce method process?

1 point

- ☐ A) All data produced by a single mapper.
- ☐ B) All the data in a single input file.
- ☐ c) All data for a given value, regardless of which mapper(s) produced it

- ☒ D) All data for a given key, regardless of which mapper(s) produced it.

Which of the following statements are true? Choose multiple answers. 1 point

- ☒ Distributed Systems can consistently give any two of the CAP properties.
- ☒ When a DataNode is down, Availability is still achieved
- ☐ When a DataNode is down. Eventual Consistency is achieved.
- ☐ When a DataNode is down, Partition Tolerance is not achieved.

Parts of Yarn ? ( select more than one) 1 point

- ☐ A ) NodeManager
- ☒ B) Scheduler and Resource manager
- ☐ C) Namenode
- ☐ D) Datanode

What happens when a DataNode goes down? Mark which false. 1 point

- ☒ All the blocks that are under-replicated get replicated to some other DataNode
- ☐ NameNode stops receiving heartbeat from this DataNode

- ☐ All the jobs scheduled on this DataNode are rescheduled to somewhere else
- ☐ NameNode deletes all the references it contains for this DataNode

A file in HDFS that is smaller than a single block size

1 point

- ☐ A) Cannot be stored in HDFS
- ☒ B) Occupies the full block's size.
- ☐ C) Occupies only the size it needs and not the full block.
- ☐ D) Can span over multiple blocks.

The need for data replication can arise in various scenarios like :

1 point

- ☐ a) Replication Factor is changed
- ☐ b) DataNode goes down
- ☐ c) Data Blocks get corrupted
- ☒ d) All of the mentioned

Which of the following is a valid flow in Hadoop ?

1 point

- ☐ a. Input -> Reducer -> Mapper -> Combiner -> -> Output
- ☐ b. Input -> Mapper -> Reducer -> Combiner -> Output
- ☒ c. Input -> Mapper -> Combiner -> Reducer -> Output
- ☐ d. Input -> Reducer -> Combiner -> Mapper -> Output

Which component is responsible for handling containers for map and reduce tasks?

1 point

- ☒ Application Manager
- ☐ Application Master

- ☐ Resource Manager
- ☐ Scheduler

Which one of these tools give better throughput and latency for seek operations?

1 point

- ☐ HDFS
- ☒ Redis
- ☐ MySql
- ☐ Postgres

During start up, the \_\_\_\_\_ loads the file system state from the fsimage and the edits log file.

1 point

- ☐ a) DataNode
- ☐ b) NameNode
- ☒ c) ActionNode
- ☐ d) None of the above

UnderReplication in HDFS means-

1 point

- ☐ A - No replication is happening in the data nodes.
- ☐ B - Replication process is very slow in the data nodes.
- ☐ C - The frequency of replication in data nodes is very low.

- ☒ D - The number of replicated copies is less than as specified by the replication factor.

In HBase, table cells are indexed by

1 point

- ☐ row key + column key + timestamp
- ☐ row key + timestamp
- ☐ row key
- ☒ None of the above

Which of the following is true about MapReduce?

1 point

- ☐ A) Data processing layer of hadoop
- ☐ B) It provides the resource management
- ☒ C) It is an open source data warehouse system for querying and analyzing large datasets stored in hadoop files
- ☐ D) All of the above

Which among the following is the duties of the DataNodes

1 point

- ☐ A) Manage file system namespace
- ☐ B) Stores meta-data
- ☐ C) Regulates client's access to files
- ☒ D) Perform read-write operation as per request for the clients

Point out the wrong statement

1 point

- ☐ Replication Factor can be configured at a cluster level (Default is set to 3) and also at a file level
- ☒ Block Report from each DataNode contains a list of all the blocks that are stored on that DataNode

- ☐ User data is stored on the local file system of DataNodes
- ☐ DataNode is aware of the files to which the blocks stored on it belong to



For a HDFS directory the replication factor(RF) is

1 point

- ☐ A - same as the RF of the files in that directory
- ☐ B - Zero
- ☐ C - 3
- ☒ D - Does not apply.

For a Spark Program, driver takes 1Gb and 1 core and two instances of executors with 4Gb and 2 cores each,determine the resources and number of tasks running at a point in time.

1 point

- ☐ a. 8Gb 4cores and 4 tasks
- ☒ b. 9Gb 5cores and 4 tasks
- ☐ c. 9Gb 5cores and 2 tasks
- ☐ d. 8Gb 5cores and 2 tasks

The \_\_\_\_\_ is a framework-specific entity that negotiates resources from the ResourceManager

1 point

- ☐ a) NodeManager
- ☒ b) ApplicationManager

☐ c) Resource Negotiator

☐ d) ApplicationMaster



I'm not a robot



reCAPTCHA

[Privacy](#) - [Terms](#)

BACK

SUBMIT

Never submit passwords through Google Forms.

This form was created inside of Media.Net. [Report Abuse](#) - [Terms of Service](#) - [Additional Terms](#)

Google Forms